

ONE HOT ENCODING

Imports libraries

```
In [1]: vocab_size = 100000 #number of unique words
        embedding_dim = 16
        max_length = 80
        trunc_type='post'
        padding_type='post'
        oov_tok = "<OOV>"
        training_size = 200000
```

<https://www.actuia.com/keras/debuter-avec-le-modele-sequentiel-de-keras/>

https://inside-machinelearning.com/en/efficient-sentences-embedding-visualization-tsne/#Preparing_the_data_-_GoEmotions

<https://vitalflux.com/keras-categorical-cross-entropy-loss-function/> sur les loss :
categorical_cross, binary_cross

```
In [2]: from sklearn.model_selection import train_test_split
        import tensorflow as tf
        from tensorflow import keras
        import tensorflow_datasets as tfds
        import pandas as pd
        import numpy as np
        %matplotlib inline
        import matplotlib as mpl
        import matplotlib.cm as cm
        import matplotlib.pyplot as plt
        import seaborn as sns
        from keras.utils import plot_model
        import os
        print(tf.__version__)
        #!pip install -q tensorflow-datasets
        from sklearn.model_selection import train_test_split
        from sklearn.datasets import make_moons
        from sklearn.ensemble import RandomForestClassifier
        from sklearn.linear_model import LogisticRegression
        from sklearn.svm import SVC
        from sklearn.ensemble import ExtraTreesClassifier
        from sklearn.ensemble import VotingClassifier
        from sklearn.metrics import accuracy_score
        from tensorflow.keras.preprocessing.text import Tokenizer
        from tensorflow.keras.preprocessing.sequence import pad_sequences
        from keras.layers import Dense, Activation, Embedding, Flatten, GlobalMaxPool1D, Dropout
        from keras.callbacks import ReduceLROnPlateau, EarlyStopping, ModelCheckpoint
```

2.9.2

Load and transform the data

We concatenate the 3 dataframes so we can deal with a single one. We drop all the useless columns in the dataframe and only keep the sentiment and the tweet associated to it.

We also reset the index because we concatenated the 3 goemotions so without that we would have 3 rows with the same index

```
In [3]: goemotions_1 = pd.read_csv('goemotions_1.csv') #Load the dataset
        goemotions_2 = pd.read_csv('goemotions_2.csv') #Load the dataset
        goemotions_3 = pd.read_csv('goemotions_3.csv') #Load the dataset
```

```
In [4]: # concat 3 df
        frames = [goemotions_1, goemotions_2, goemotions_3]
        df = pd.concat(frames)
        # drop others columns
        df.drop(['id', 'author', 'subreddit', 'link_id', 'parent_id', 'created_utc', 'rater_id', ''],
               inplace=True)
        df.reset_index(drop=True, inplace=True)
        df.shape
```

Out[4]: (211225, 29)

We assign the following variables :


- sentences : the column "text" of df that refers to all the tweets in the dataframe
- labels : sentiment get dummies of the dataframe, all the sentiment columns so from the index 1 (index 0 being the text) to the last column) There are 28 sentiments in the dataframe

```
In [5]: sentences = df['text']
        # df of all the 28 emotions
        labels = df.iloc[:, 1:]
        #Labels.shape
        labels
```

```
Out[5]:
```

| | admiration | amusement | anger | annoyance | approval | caring | confusion | curiosity | desire |
|--------|------------|-----------|-------|-----------|----------|--------|-----------|-----------|--------|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 2 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 4 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 211220 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 211221 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 |
| 211222 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 211223 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 |
| 211224 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

211225 rows × 28 columns



Create train and test

We associate to training_sentences the tweets for the training size so from 0 to 200 000

The testing_sentences is all the tweet from 200 000 to the end of the dataframe

The training_labels takes all the labels associated to the tweets for the training size (0 to 200 000) The testing_labels takes the remaining labels

```
In [6]: # training_sentences select the text, testing_labels the labels associated
training_sentences = sentences[0:training_size]
testing_sentences = sentences[training_size:]
training_labels = labels[0:training_size]
testing_labels = labels[training_size:]

# training_sentences[0] is the first tweet in training_sentences
```

We tokenize training_sentences and testing_sentences

```
In [7]: tokenizer = Tokenizer(num_words=vocab_size, oov_token=oov_tok)
tokenizer.fit_on_texts(training_sentences)

word_index = tokenizer.word_index

# make all the sentences of equal size, add zeros at the end of the sentences to fi

training_sequences = tokenizer.texts_to_sequences(training_sentences)
training_padded = pad_sequences(training_sequences, maxlen=max_length, padding=paddi

testing_sequences = tokenizer.texts_to_sequences(testing_sentences)
testing_padded = pad_sequences(testing_sequences, maxlen=max_length, padding=padding
```

```
In [8]: # Need this block to get it to work with TensorFlow 2.x
training_padded = np.array(training_padded)
training_labels = np.array(training_labels)
testing_padded = np.array(testing_padded)
testing_labels = np.array(testing_labels)
```

ACURACY METRIC

MODEL 1 loss = binary cross entropy

```
In [ ]:
```

Hyperparameters

The whole dataset is 211225 so we split inot a training size of 200 000 and the rest is the test set.

We set the maximum vocabulary size to 10 000 and the maximum length of a tweet to 150

703 is the longest tweet size

```
In [18]: model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),
    tf.keras.layers.GlobalAveragePooling1D(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(28, activation='sigmoid')
])
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
model.summary()
```

Model: "sequential_4"

| Layer (type) | Output Shape | Param # |
|--|----------------|---------|
| ===== | | |
| embedding_4 (Embedding) | (None, 80, 16) | 1600000 |
| global_average_pooling1d_4 (GlobalAveragePooling1D) | (None, 16) | 0 |
| dense_8 (Dense) | (None, 128) | 2176 |
| dense_9 (Dense) | (None, 28) | 3612 |
| ===== | | |
| Total params: 1,605,788 | | |
| Trainable params: 1,605,788 | | |
| Non-trainable params: 0 | | |

```
In [19]: num_epochs = 20
history = model.fit(training_padded, training_labels, epochs=num_epochs, validation_
```

```
Epoch 1/20
6250/6250 - 17s - loss: 0.1598 - accuracy: 0.2652 - val_loss: 0.1478 - val_accuracy:
0.2971 - 17s/epoch - 3ms/step
Epoch 2/20
6250/6250 - 16s - loss: 0.1425 - accuracy: 0.3158 - val_loss: 0.1369 - val_accuracy:
0.3395 - 16s/epoch - 3ms/step
Epoch 3/20
6250/6250 - 16s - loss: 0.1311 - accuracy: 0.3650 - val_loss: 0.1299 - val_accuracy:
0.3668 - 16s/epoch - 3ms/step
Epoch 4/20
6250/6250 - 16s - loss: 0.1239 - accuracy: 0.3881 - val_loss: 0.1258 - val_accuracy:
0.3840 - 16s/epoch - 3ms/step
Epoch 5/20
6250/6250 - 17s - loss: 0.1191 - accuracy: 0.4063 - val_loss: 0.1245 - val_accuracy:
0.3903 - 17s/epoch - 3ms/step
Epoch 6/20
6250/6250 - 16s - loss: 0.1157 - accuracy: 0.4182 - val_loss: 0.1236 - val_accuracy:
0.3955 - 16s/epoch - 3ms/step
Epoch 7/20
6250/6250 - 16s - loss: 0.1132 - accuracy: 0.4281 - val_loss: 0.1228 - val_accuracy:
0.3919 - 16s/epoch - 3ms/step
Epoch 8/20
6250/6250 - 16s - loss: 0.1111 - accuracy: 0.4355 - val_loss: 0.1232 - val_accuracy:
```

```

0.3943 - 16s/epoch - 3ms/step
Epoch 9/20
6250/6250 - 17s - loss: 0.1093 - accuracy: 0.4425 - val_loss: 0.1228 - val_accuracy:
0.3953 - 17s/epoch - 3ms/step
Epoch 10/20
6250/6250 - 16s - loss: 0.1076 - accuracy: 0.4487 - val_loss: 0.1229 - val_accuracy:
0.3926 - 16s/epoch - 3ms/step
Epoch 11/20
6250/6250 - 16s - loss: 0.1060 - accuracy: 0.4537 - val_loss: 0.1236 - val_accuracy:
0.3939 - 16s/epoch - 3ms/step
Epoch 12/20
6250/6250 - 16s - loss: 0.1047 - accuracy: 0.4576 - val_loss: 0.1241 - val_accuracy:
0.3962 - 16s/epoch - 3ms/step
Epoch 13/20
6250/6250 - 16s - loss: 0.1034 - accuracy: 0.4612 - val_loss: 0.1244 - val_accuracy:
0.3887 - 16s/epoch - 3ms/step
Epoch 14/20
6250/6250 - 16s - loss: 0.1022 - accuracy: 0.4652 - val_loss: 0.1254 - val_accuracy:
0.3861 - 16s/epoch - 3ms/step
Epoch 15/20
6250/6250 - 16s - loss: 0.1012 - accuracy: 0.4678 - val_loss: 0.1250 - val_accuracy:
0.3968 - 16s/epoch - 3ms/step
Epoch 16/20
6250/6250 - 16s - loss: 0.1002 - accuracy: 0.4702 - val_loss: 0.1255 - val_accuracy:
0.3935 - 16s/epoch - 3ms/step
Epoch 17/20
6250/6250 - 17s - loss: 0.0993 - accuracy: 0.4732 - val_loss: 0.1265 - val_accuracy:
0.3960 - 17s/epoch - 3ms/step
Epoch 18/20
6250/6250 - 20s - loss: 0.0985 - accuracy: 0.4744 - val_loss: 0.1273 - val_accuracy:
0.3903 - 20s/epoch - 3ms/step
Epoch 19/20
6250/6250 - 18s - loss: 0.0977 - accuracy: 0.4772 - val_loss: 0.1282 - val_accuracy:
0.3751 - 18s/epoch - 3ms/step
Epoch 20/20
6250/6250 - 18s - loss: 0.0970 - accuracy: 0.4796 - val_loss: 0.1291 - val_accuracy:
0.3906 - 18s/epoch - 3ms/step

```

In [20]:


```

import matplotlib.pyplot as plt

def plot_graphs(history, string):
    plt.plot(history.history[string])
    plt.plot(history.history['val_'+string])
    plt.xlabel("Epochs")
    plt.ylabel(string)
    plt.legend([string, 'val_'+string])
    plt.show()

plot_graphs(history, "accuracy")
plot_graphs(history, "loss")

```



Conclusion :

the model seems to behave well after each iteration of optimization but the loss stagnates and so the accuracy doesn't go beyond 0.4. The model tends to overfit a little bit at the end.

I try another way to predict the sentiment by having a label encoding instead of a one hot encoding process

MODEL 1 WITH DIFFERENT PARAMETERS

vocab_size

AUC METRIC

MODEL 1

```
In [16]: model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),
    tf.keras.layers.GlobalAveragePooling1D(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(28, activation='sigmoid')
```

```

])
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=[tf.keras.metrics
model.summary()

num_epochs = 20
history = model.fit(training_padded, training_labels, epochs=num_epochs, validation_

```

Model: "sequential_3"

| Layer (type) | Output Shape | Param # |
|--|----------------|---------|
| embedding_3 (Embedding) | (None, 80, 16) | 1600000 |
| global_average_pooling1d_3 (GlobalAveragePooling1D) | (None, 16) | 0 |
| dense_6 (Dense) | (None, 128) | 2176 |
| dense_7 (Dense) | (None, 28) | 3612 |
| Total params: 1,605,788 | | |
| Trainable params: 1,605,788 | | |
| Non-trainable params: 0 | | |

Epoch 1/20

6250/6250 - 21s - loss: 0.1610 - auc_2: 0.7355 - val_loss: 0.1503 - val_auc_2: 0.7714 - 21s/epoch - 3ms/step

Epoch 2/20

6250/6250 - 21s - loss: 0.1437 - auc_2: 0.8051 - val_loss: 0.1380 - val_auc_2: 0.8267 - 21s/epoch - 3ms/step

Epoch 3/20

6250/6250 - 20s - loss: 0.1320 - auc_2: 0.8488 - val_loss: 0.1309 - val_auc_2: 0.8526 - 20s/epoch - 3ms/step

Epoch 4/20

6250/6250 - 20s - loss: 0.1252 - auc_2: 0.8724 - val_loss: 0.1269 - val_auc_2: 0.8676 - 20s/epoch - 3ms/step

Epoch 5/20

6250/6250 - 20s - loss: 0.1201 - auc_2: 0.8882 - val_loss: 0.1243 - val_auc_2: 0.8760 - 20s/epoch - 3ms/step

Epoch 6/20

6250/6250 - 20s - loss: 0.1162 - auc_2: 0.8984 - val_loss: 0.1232 - val_auc_2: 0.8789 - 20s/epoch - 3ms/step

Epoch 7/20

6250/6250 - 20s - loss: 0.1134 - auc_2: 0.9051 - val_loss: 0.1226 - val_auc_2: 0.8808 - 20s/epoch - 3ms/step

Epoch 8/20

6250/6250 - 20s - loss: 0.1111 - auc_2: 0.9105 - val_loss: 0.1228 - val_auc_2: 0.8803 - 20s/epoch - 3ms/step

Epoch 9/20

6250/6250 - 20s - loss: 0.1090 - auc_2: 0.9151 - val_loss: 0.1231 - val_auc_2: 0.8812 - 20s/epoch - 3ms/step

Epoch 10/20

6250/6250 - 20s - loss: 0.1073 - auc_2: 0.9191 - val_loss: 0.1232 - val_auc_2: 0.8806 - 20s/epoch - 3ms/step

Epoch 11/20

6250/6250 - 20s - loss: 0.1057 - auc_2: 0.9225 - val_loss: 0.1238 - val_auc_2: 0.8794 - 20s/epoch - 3ms/step

Epoch 12/20

6250/6250 - 20s - loss: 0.1043 - auc_2: 0.9253 - val_loss: 0.1243 - val_auc_2: 0.8797 - 20s/epoch - 3ms/step

Epoch 13/20

6250/6250 - 20s - loss: 0.1030 - auc_2: 0.9280 - val_loss: 0.1249 - val_auc_2: 0.8778 - 20s/epoch - 3ms/step

Epoch 14/20

6250/6250 - 20s - loss: 0.1019 - auc_2: 0.9304 - val_loss: 0.1255 - val_auc_2: 0.877

3 - 20s/epoch - 3ms/step

Epoch 15/20

6250/6250 - 20s - loss: 0.1009 - auc_2: 0.9323 - val_loss: 0.1259 - val_auc_2: 0.876

7 - 20s/epoch - 3ms/step

Epoch 16/20

6250/6250 - 20s - loss: 0.1000 - auc_2: 0.9342 - val_loss: 0.1267 - val_auc_2: 0.875

7 - 20s/epoch - 3ms/step

Epoch 17/20

6250/6250 - 21s - loss: 0.0992 - auc_2: 0.9358 - val_loss: 0.1273 - val_auc_2: 0.874

8 - 21s/epoch - 3ms/step

Epoch 18/20

6250/6250 - 21s - loss: 0.0984 - auc_2: 0.9374 - val_loss: 0.1289 - val_auc_2: 0.873

3 - 21s/epoch - 3ms/step

Epoch 19/20

6250/6250 - 20s - loss: 0.0977 - auc_2: 0.9387 - val_loss: 0.1291 - val_auc_2: 0.870

9 - 20s/epoch - 3ms/step

Epoch 20/20

6250/6250 - 20s - loss: 0.0971 - auc_2: 0.9400 - val_loss: 0.1302 - val_auc_2: 0.870


4 - 20s/epoch - 3ms/step


In [17]:

```
import matplotlib.pyplot as plt

def plot_graphs(history, string):
    plt.plot(history.history[string])
    plt.plot(history.history['val_'+string])
    plt.xlabel("Epochs")
    plt.ylabel(string)
    plt.legend([string, 'val_'+string])
    plt.show()

plot_graphs(history, "auc_2")
plot_graphs(history, "loss")
```





The loss decreases and the accuracy increases so the training process is efficient. There is no overfitting

MODEL 1 with dropout

```
In [14]: model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.GlobalAveragePooling1D(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(28, activation='sigmoid')
])
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=[tf.keras.metrics
model.summary()

num_epochs = 20
history = model.fit(training_padded, training_labels, epochs=num_epochs, validation_
```

Model: "sequential_2"

| Layer (type) | Output Shape | Param # |
|--|----------------|---------|
| ===== | | |
| embedding_2 (Embedding) | (None, 80, 16) | 1600000 |
| dropout_1 (Dropout) | (None, 80, 16) | 0 |
| global_average_pooling1d_2 (GlobalAveragePooling1D) | (None, 16) | 0 |
| dense_4 (Dense) | (None, 128) | 2176 |
| dense_5 (Dense) | (None, 28) | 3612 |

=====

Total params: 1,605,788

Trainable params: 1,605,788

Non-trainable params: 0

Epoch 1/20

6250/6250 - 23s - loss: 0.1604 - auc_1: 0.7370 - val_loss: 0.1491 - val_auc_1: 0.777
4 - 23s/epoch - 4ms/step

Epoch 2/20

6250/6250 - 21s - loss: 0.1458 - auc_1: 0.7954 - val_loss: 0.1401 - val_auc_1: 0.816
9 - 21s/epoch - 3ms/step

```

Epoch 3/20
6250/6250 - 21s - loss: 0.1381 - auc_1: 0.8273 - val_loss: 0.1335 - val_auc_1: 0.842
7 - 21s/epoch - 3ms/step
Epoch 4/20
6250/6250 - 20s - loss: 0.1326 - auc_1: 0.8478 - val_loss: 0.1294 - val_auc_1: 0.858
2 - 20s/epoch - 3ms/step
Epoch 5/20
6250/6250 - 20s - loss: 0.1293 - auc_1: 0.8595 - val_loss: 0.1274 - val_auc_1: 0.866
1 - 20s/epoch - 3ms/step
Epoch 6/20
6250/6250 - 21s - loss: 0.1273 - auc_1: 0.8663 - val_loss: 0.1260 - val_auc_1: 0.870
2 - 21s/epoch - 3ms/step
Epoch 7/20
6250/6250 - 22s - loss: 0.1258 - auc_1: 0.8714 - val_loss: 0.1253 - val_auc_1: 0.873
4 - 22s/epoch - 3ms/step
Epoch 8/20
6250/6250 - 21s - loss: 0.1245 - auc_1: 0.8752 - val_loss: 0.1247 - val_auc_1: 0.874
5 - 21s/epoch - 3ms/step
Epoch 9/20
6250/6250 - 20s - loss: 0.1236 - auc_1: 0.8777 - val_loss: 0.1248 - val_auc_1: 0.873
1 - 20s/epoch - 3ms/step
Epoch 10/20
6250/6250 - 20s - loss: 0.1226 - auc_1: 0.8807 - val_loss: 0.1237 - val_auc_1: 0.877
2 - 20s/epoch - 3ms/step
Epoch 11/20
6250/6250 - 21s - loss: 0.1218 - auc_1: 0.8828 - val_loss: 0.1232 - val_auc_1: 0.879
1 - 21s/epoch - 3ms/step
Epoch 12/20
6250/6250 - 20s - loss: 0.1211 - auc_1: 0.8850 - val_loss: 0.1228 - val_auc_1: 0.879
6 - 20s/epoch - 3ms/step
Epoch 13/20
6250/6250 - 20s - loss: 0.1205 - auc_1: 0.8865 - val_loss: 0.1226 - val_auc_1: 0.880
6 - 20s/epoch - 3ms/step
Epoch 14/20
6250/6250 - 20s - loss: 0.1199 - auc_1: 0.8879 - val_loss: 0.1226 - val_auc_1: 0.880
1 - 20s/epoch - 3ms/step
Epoch 15/20
6250/6250 - 21s - loss: 0.1194 - auc_1: 0.8894 - val_loss: 0.1222 - val_auc_1: 0.882
0 - 21s/epoch - 3ms/step
Epoch 16/20
6250/6250 - 21s - loss: 0.1190 - auc_1: 0.8903 - val_loss: 0.1224 - val_auc_1: 0.880
9 - 21s/epoch - 3ms/step
Epoch 17/20
6250/6250 - 20s - loss: 0.1184 - auc_1: 0.8919 - val_loss: 0.1225 - val_auc_1: 0.880
1 - 20s/epoch - 3ms/step
Epoch 18/20
6250/6250 - 21s - loss: 0.1181 - auc_1: 0.8928 - val_loss: 0.1227 - val_auc_1: 0.878
9 - 21s/epoch - 3ms/step
Epoch 19/20
6250/6250 - 21s - loss: 0.1177 - auc_1: 0.8935 - val_loss: 0.1220 - val_auc_1: 0.882
7 - 21s/epoch - 3ms/step
Epoch 20/20
6250/6250 - 21s - loss: 0.1174 - auc_1: 0.8944 - val_loss: 0.1223 - val_auc_1: 0.881
3 - 21s/epoch - 3ms/step

```

In [15]:

```


import matplotlib.pyplot as plt

def plot_graphs(history, string):
    plt.plot(history.history[string])
    plt.plot(history.history['val_'+string])
    plt.xlabel("Epochs")
    plt.ylabel(string)
    plt.legend([string, 'val_'+string])

```

```
plt.show()

plot_graphs(history, "auc_1")
plot_graphs(history, "loss")
```



Other model evaluated less conclusive

MODEL 2

```
In [ ]: from keras.layers import LSTM
        from keras.layers import Bidirectional

        model = tf.keras.Sequential([
            tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),
            tf.keras.layers.Bidirectional(LSTM(8, return_sequences=True)),
            tf.keras.layers.GlobalAveragePooling1D(),
            tf.keras.layers.Dense(24, activation='relu'),
            tf.keras.layers.Dense(28, activation='sigmoid') # change its shape to the maxi
        ])

        model.compile(loss='binary_crossentropy', optimizer='adam', metrics=[tf.keras.metrics.

        #model.summary()
```

loss='categorical_crossentropy' : auc_2: 0.5032

```
In [ ]: # keras.utils.plot_model(model, "27 sentiments analysis.png", show_shapes=True)
```

```
In [ ]: from keras.callbacks import LearningRateScheduler

# # we add Learning rate scheduler to change because we are on a plateau
# # for that we add it in callback

# callbacks = [
#     ReduceLROnPlateau(), # monitors a quantity and if no improvement is seen for a
#     ModelCheckpoint(filepath='model_sparse', save_best_only=True)
# ]



#####
def lr_scheduler(epoch, lr):
    decay_rate = 0.1
    decay_step = 90
    if epoch % decay_step == 0 and epoch:
        return lr * decay_rate
    return lr
callbacks = [
    keras.callbacks.LearningRateScheduler(lr_scheduler, verbose=1)
]
#####
#callbacks = keras.callbacks.ModelCheckpoint("model_sparse.h5", save_best_only=True)
```

```
In [ ]: num_epochs = 10
history = model.fit(training_padded, training_labels, epochs=num_epochs, validation_
```

```
Epoch 1/10
6250/6250 - 60s - loss: 0.1570 - auc_3: 0.7611 - val_loss: 0.1443 - val_auc_3: 0.801
5 - 60s/epoch - 10ms/step
Epoch 2/10
6250/6250 - 57s - loss: 0.1412 - auc_3: 0.8171 - val_loss: 0.1377 - val_auc_3: 0.828
9 - 57s/epoch - 9ms/step
Epoch 3/10
6250/6250 - 57s - loss: 0.1340 - auc_3: 0.8440 - val_loss: 0.1341 - val_auc_3: 0.840
4 - 57s/epoch - 9ms/step
Epoch 4/10
6250/6250 - 56s - loss: 0.1293 - auc_3: 0.8585 - val_loss: 0.1312 - val_auc_3: 0.851
9 - 56s/epoch - 9ms/step
Epoch 5/10
6250/6250 - 56s - loss: 0.1244 - auc_3: 0.8729 - val_loss: 0.1288 - val_auc_3: 0.859
6 - 56s/epoch - 9ms/step
Epoch 6/10
6250/6250 - 57s - loss: 0.1203 - auc_3: 0.8844 - val_loss: 0.1271 - val_auc_3: 0.866
0 - 57s/epoch - 9ms/step
Epoch 7/10
6250/6250 - 56s - loss: 0.1172 - auc_3: 0.8923 - val_loss: 0.1267 - val_auc_3: 0.867
2 - 56s/epoch - 9ms/step
Epoch 8/10
6250/6250 - 56s - loss: 0.1149 - auc_3: 0.8982 - val_loss: 0.1261 - val_auc_3: 0.870
0 - 56s/epoch - 9ms/step
Epoch 9/10
6250/6250 - 56s - loss: 0.1129 - auc_3: 0.9032 - val_loss: 0.1262 - val_auc_3: 0.870
2 - 56s/epoch - 9ms/step
Epoch 10/10
6250/6250 - 56s - loss: 0.1112 - auc_3: 0.9070 - val_loss: 0.1268 - val_auc_3: 0.868
6 - 56s/epoch - 9ms/step
```

```
In [ ]: def plot_graphs(history, string):
plt.plot(history.history[string])
plt.plot(history.history['val_'+string])
plt.xlabel("Epochs")
plt.ylabel(string)
plt.legend([string, 'val_'+string])
plt.show()

plot_graphs(history, "auc_3")
plot_graphs(history, "loss")
```



MODEL 3

```
In [ ]: # Here the Length refers to the number of rows of the filter,
# here it is the dimension of the entire word embedding or the entire character repr
filter_length = 300

model = tf.keras.Sequential()
model.add(Embedding(vocab_size, 20, input_length=max_length))
#model.add(Dropout(0.5))
model.add(Conv1D(filter_length, 3, padding='valid', activation='relu', strides=1))
model.add(GlobalMaxPool1D())
model.add(Dense(training_labels.shape[1]))
model.add(Activation('sigmoid'))

#model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=[tf.keras.metric

callbacks = [
    ReduceLROnPlateau(),
    ModelCheckpoint(filepath='model-conv1d.h5', save_best_only=True)
]

num_epochs= 10
history = model.fit(training_padded, training_labels,
                    epochs=num_epochs, callbacks=callbacks, validation_data=(testing

```

```

Epoch 1/10
6250/6250 - 26s - loss: 0.1321 - auc_4: 0.8532 - val_loss: 0.1205 - val_auc_4: 0.887
4 - lr: 0.0010 - 26s/epoch - 4ms/step
Epoch 2/10
6250/6250 - 22s - loss: 0.1158 - auc_4: 0.9000 - val_loss: 0.1180 - val_auc_4: 0.894
5 - lr: 0.0010 - 22s/epoch - 3ms/step
Epoch 3/10
6250/6250 - 22s - loss: 0.1105 - auc_4: 0.9134 - val_loss: 0.1177 - val_auc_4: 0.895
5 - lr: 0.0010 - 22s/epoch - 3ms/step
Epoch 4/10
6250/6250 - 23s - loss: 0.1066 - auc_4: 0.9218 - val_loss: 0.1183 - val_auc_4: 0.895
0 - lr: 0.0010 - 23s/epoch - 4ms/step
Epoch 5/10
6250/6250 - 21s - loss: 0.1037 - auc_4: 0.9284 - val_loss: 0.1199 - val_auc_4: 0.892
2 - lr: 0.0010 - 21s/epoch - 3ms/step
Epoch 6/10
6250/6250 - 21s - loss: 0.1014 - auc_4: 0.9332 - val_loss: 0.1210 - val_auc_4: 0.889
8 - lr: 0.0010 - 21s/epoch - 3ms/step
Epoch 7/10
6250/6250 - 21s - loss: 0.0995 - auc_4: 0.9370 - val_loss: 0.1226 - val_auc_4: 0.888
1 - lr: 0.0010 - 21s/epoch - 3ms/step
Epoch 8/10
6250/6250 - 21s - loss: 0.0980 - auc_4: 0.9402 - val_loss: 0.1238 - val_auc_4: 0.884
5 - lr: 0.0010 - 21s/epoch - 3ms/step
Epoch 9/10
6250/6250 - 20s - loss: 0.0967 - auc_4: 0.9428 - val_loss: 0.1256 - val_auc_4: 0.881
3 - lr: 0.0010 - 20s/epoch - 3ms/step
Epoch 10/10
6250/6250 - 21s - loss: 0.0956 - auc_4: 0.9450 - val_loss: 0.1261 - val_auc_4: 0.882
4 - lr: 0.0010 - 21s/epoch - 3ms/step

```



```

In [ ]: import matplotlib.pyplot as plt

def plot_graphs(history, string):
    plt.plot(history.history[string])
    plt.plot(history.history['val_'+string])
    plt.xlabel("Epochs")
    plt.ylabel(string)
    plt.legend([string, 'val_'+string])
    plt.show()

plot_graphs(history, "auc_4")
plot_graphs(history, "loss")

```



MODEL 3

sans callback

avec dropout

<https://stackabuse.com/python-for-nlp-multi-label-text-classification-with-keras/>

<https://www.kaggle.com/code/bansodesandeep/multilabel-cuisine-classification-cnn-dnn-lstm>

```
In [ ]: # Here the length refers to the number of rows of the filter,
# here it is the dimension of the entire word embedding or the entire character repr
filter_length = 300

model = tf.keras.Sequential()
model.add(Embedding(vocab_size, 20, input_length=max_length))
model.add(Dropout(0.5))
model.add(Conv1D(filter_length, 3, padding='valid', activation='relu', strides=1))
model.add(GlobalMaxPool1D())
model.add(Dense(training_labels.shape[1]))
model.add(Activation('sigmoid'))

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=[tf.keras.metric
```

```

callbacks = [
    ReduceLROnPlateau(),
    ModelCheckpoint(filepath='model-conv1d.h5', save_best_only=True)
]

num_epochs= 10
history = model.fit(training_padded, training_labels,
                    epochs=num_epochs, validation_data=(testing_padded, testing_labels))

```

```

Epoch 1/10
6250/6250 - 25s - loss: 0.1407 - auc_7: 0.7028 - val_loss: 0.1479 - val_auc_7: 0.797
9 - 25s/epoch - 4ms/step
Epoch 2/10
6250/6250 - 23s - loss: 0.1284 - auc_7: 0.7878 - val_loss: 0.1603 - val_auc_7: 0.815
6 - 23s/epoch - 4ms/step
Epoch 3/10
6250/6250 - 24s - loss: 0.1252 - auc_7: 0.8070 - val_loss: 0.1807 - val_auc_7: 0.812
7 - 24s/epoch - 4ms/step
Epoch 4/10
6250/6250 - 26s - loss: 0.1231 - auc_7: 0.8166 - val_loss: 0.1927 - val_auc_7: 0.816
4 - 26s/epoch - 4ms/step
Epoch 5/10
6250/6250 - 23s - loss: 0.1216 - auc_7: 0.8246 - val_loss: 0.2038 - val_auc_7: 0.817
4 - 23s/epoch - 4ms/step
Epoch 6/10
6250/6250 - 24s - loss: 0.1202 - auc_7: 0.8295 - val_loss: 0.2154 - val_auc_7: 0.822
1 - 24s/epoch - 4ms/step
Epoch 7/10
6250/6250 - 23s - loss: 0.1193 - auc_7: 0.8339 - val_loss: 0.2172 - val_auc_7: 0.826
0 - 23s/epoch - 4ms/step
Epoch 8/10
6250/6250 - 23s - loss: 0.1182 - auc_7: 0.8388 - val_loss: 0.2206 - val_auc_7: 0.825
2 - 23s/epoch - 4ms/step
Epoch 9/10
6250/6250 - 23s - loss: 0.1176 - auc_7: 0.8413 - val_loss: 0.2287 - val_auc_7: 0.831
0 - 23s/epoch - 4ms/step
Epoch 10/10
6250/6250 - 23s - loss: 0.1169 - auc_7: 0.8444 - val_loss: 0.2413 - val_auc_7: 0.836
2 - 23s/epoch - 4ms/step

```

In []:



```

import matplotlib.pyplot as plt

def plot_graphs(history, string):
    plt.plot(history.history[string])
    plt.plot(history.history['val_'+string])
    plt.xlabel("Epochs")
    plt.ylabel(string)
    plt.legend([string, 'val_'+string])
    plt.show()

plot_graphs(history, "auc_7")
plot_graphs(history, "loss")

```

MODEL 3

sans callback

avec dropout

sans filter length, kernel size null

```
In [ ]: # Here the length refers to the number of rows of the filter,
# here it is the dimension of the entire word embedding or the entire character repr
filter_length = 100

model = tf.keras.Sequential()
model.add(Embedding(vocab_size, 20, input_length=max_length))
model.add(Dropout(0.5))
model.add(Conv1D(filter_length, 3, padding='valid', activation='relu', strides=1))
model.add(GlobalMaxPool1D())
model.add(Dense(training_labels.shape[1]))
model.add(Activation('sigmoid'))

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=[tf.keras.metric
```

```

callbacks = [
    ReduceLROnPlateau(),
    ModelCheckpoint(filepath='model-conv1d.h5', save_best_only=True)
]

num_epochs= 10
history = model.fit(training_padded, training_labels,
                    epochs=num_epochs, validation_data=(testing_padded, testing_labels))

```

```

Epoch 1/10
6250/6250 - 26s - loss: 0.1445 - auc_10: 0.6742 - val_loss: 0.1519 - val_auc_10: 0.7753 - 26s/epoch - 4ms/step
Epoch 2/10
6250/6250 - 23s - loss: 0.1314 - auc_10: 0.7669 - val_loss: 0.1537 - val_auc_10: 0.8055 - 23s/epoch - 4ms/step
Epoch 3/10
6250/6250 - 24s - loss: 0.1278 - auc_10: 0.7920 - val_loss: 0.1585 - val_auc_10: 0.8093 - 24s/epoch - 4ms/step
Epoch 4/10
6250/6250 - 23s - loss: 0.1254 - auc_10: 0.8050 - val_loss: 0.1699 - val_auc_10: 0.8068 - 23s/epoch - 4ms/step
Epoch 5/10
6250/6250 - 23s - loss: 0.1238 - auc_10: 0.8126 - val_loss: 0.1812 - val_auc_10: 0.8006 - 23s/epoch - 4ms/step
Epoch 6/10
6250/6250 - 23s - loss: 0.1224 - auc_10: 0.8200 - val_loss: 0.1959 - val_auc_10: 0.8048 - 23s/epoch - 4ms/step
Epoch 7/10
6250/6250 - 23s - loss: 0.1213 - auc_10: 0.8243 - val_loss: 0.2129 - val_auc_10: 0.8027 - 23s/epoch - 4ms/step
Epoch 8/10
6250/6250 - 23s - loss: 0.1205 - auc_10: 0.8279 - val_loss: 0.2172 - val_auc_10: 0.8026 - 23s/epoch - 4ms/step
Epoch 9/10
6250/6250 - 23s - loss: 0.1197 - auc_10: 0.8313 - val_loss: 0.2312 - val_auc_10: 0.7978 - 23s/epoch - 4ms/step
Epoch 10/10
6250/6250 - 23s - loss: 0.1190 - auc_10: 0.8340 - val_loss: 0.2364 - val_auc_10: 0.8024 - 23s/epoch - 4ms/step

```

In []:

```

import matplotlib.pyplot as plt

def plot_graphs(history, string):
    plt.plot(history.history[string])
    plt.plot(history.history['val_'+string])
    plt.xlabel("Epochs")
    plt.ylabel(string)
    plt.legend([string, 'val_'+string])
    plt.show()

plot_graphs(history, "auc_10")
plot_graphs(history, "loss")

```

